

CNN Org Mode

Jesus Sierralaya

April 10, 2024

Contents

| | | |
|----------|--|----------|
| 1 | Load libraries | 1 |
| 2 | Set parameters | 1 |
| 3 | Pre-process | 2 |
| 4 | Create the model | 2 |
| 5 | Set the parameters, compile and train the model | 3 |
| 6 | Evaluate the performance of the trained model | 3 |

1 Load libraries

```
import os; os.environ['TF_CPP_MIN_LOG_LEVEL'] = '2'
import matplotlib.pyplot as plt
import numpy as np
import keras
from keras import layers
import tensorflow as tf
import argparse
from sklearn.preprocessing import LabelBinarizer
from keras.datasets import mnist
```

2 Set parameters

```
num_classes = 10
input_shape = (28, 28, 1)
```

3 Pre-process

```
(x_train, y_train), (x_test, y_test) = keras.datasets.mnist.load_data()

x_train = x_train.astype("float32") / 255
x_test = x_test.astype("float32") / 255

x_train = np.expand_dims(x_train, -1)
x_test = np.expand_dims(x_test, -1)
print("x_train shape:", x_train.shape)
print(x_train.shape[0], "train samples")
print(x_test.shape[0], "test samples")

label_binarizer = LabelBinarizer()
y_train = label_binarizer.fit_transform(y_train)
y_test = label_binarizer.fit_transform(y_test)

x_train shape: (60000, 28, 28, 1)
60000 train samples
10000 test samples
```

4 Create the model

```
model = keras.Sequential(
    [
        keras.Input(shape=input_shape),
        layers.Conv2D(32, kernel_size=(3, 3), activation="relu"),
        layers.MaxPooling2D(pool_size=(2, 2)),
        layers.Conv2D(64, kernel_size=(3, 3), activation="relu"),
        layers.MaxPooling2D(pool_size=(2, 2)),
        layers.Flatten(),
        layers.Dropout(0.5),
        layers.Dense(num_classes, activation="softmax"),
    ]
)

model.summary()
```

Model: "sequential_1"

| Layer (type) | Output Shape | Param # |
|-------------------------------------|--------------------|---------|
| conv2d_2 (Conv2D) | (None, 26, 26, 32) | 320 |
| max_pooling2d_2 (MaxPooling2D) | (None, 13, 13, 32) | 0 |
| conv2d_3 (Conv2D) | (None, 11, 11, 64) | 18496 |
| max_pooling2d_3 (MaxPooling2D) | (None, 5, 5, 64) | 0 |
| flatten_1 (Flatten) | (None, 1600) | 0 |
| dropout_1 (Dropout) | (None, 1600) | 0 |
| dense_1 (Dense) | (None, 10) | 16010 |
| Total params: 34826 (136.04 KB) | | |
| Trainable params: 34826 (136.04 KB) | | |
| Non-trainable params: 0 (0.00 Byte) | | |

5 Set the parameters, compile and train the model

```
batch_size = 128
epochs = 1
```

```
model.compile(loss="categorical_crossentropy", optimizer="adam", metrics=["accuracy"])
```

```
model.fit(x_train, y_train, batch_size=batch_size, epochs=epochs, validation_split=0.1)
```

6 Evaluate the performance of the trained model

```
test_loss, test_accuracy = model.evaluate(x_test, y_test, verbose=2)
print(f"Test accuracy: {test_accuracy * 100:.2f}%")
```

```
f"{test_accuracy * 100:.2f}%"
```

The test accuracy is 97.37%.